

# Supporting Systematic Literature Reviews Using Deep-Learning-Based Language Models

Rand Alchokr  
Otto-von-Guericke University  
Magdeburg, Germany  
rand.alchokr@ovgu.de

Manoj Borkar  
Otto-von-Guericke University  
Magdeburg, Germany  
manoj.borkar@st.ovgu.de

Sharanya Thotadarya  
Otto-von-Guericke University  
Magdeburg, Germany  
sharanya.hunasamaranahalli@ovgu.de

Gunter Saake  
Otto-von-Guericke University  
Magdeburg, Germany  
saake@ovgu.de

Thomas Leich  
Harz University & METOP GmbH  
Weringerode & Magdeburg, Germany  
tleich@hs-harz.de

## ABSTRACT

**Background:** Systematic Literature Reviews are an important research method for gathering and evaluating the available evidence regarding a specific research topic. However, the process of conducting an Systematic Literature Review manually can be difficult and time-consuming. For this reason, researchers aim to semi-automate this process or some of its phases. Still, automating Systematic Literature Reviews requires more research and poses several open challenges. **Aim:** We aimed at using a deep-learning-based clustering model involving transformer-based language models and a weighted scheme to accelerate the conduct phase of Systematic Literature Reviews. In particular, we intend to support researchers with automation for efficiently scanning the initial set of retrieved publications and selecting relevant primary studies. **Method:** We performed an experiment using two manually conducted SLRs that comprise a high number of initial results (153 and 174, respectively) to evaluate the performance of two deep-learning-based clustering models. These models build on transformer-based language models (i.e., BERT and S-BERT) to extract contextualized embeddings on different text levels along with a weighted scheme to cluster similar publications. **Results:** Our primary results show that clustering based on embedding at paragraph-level using S-BERT-paragraph and weight schemes represents the best performing model setting in terms of correctly identifying primary studies that were selected in the original Systematic Literature Reviews (12 and 9, respectively). **Conclusions:** The findings indicate that using natural-language-based deep-learning architectures for semi-automating the selection of primary studies can accelerate the scanning process. While our results represent first insights only, such a technique seems to enhance the conduct phase of the SLR process, promising to help researchers identify the most relevant publications more quickly and efficiently.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

JCDL 2022, June 20–24, 2022, Cologne, Germany

© 2022 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/10.1145/000000.000000>

## CCS CONCEPTS

• General and reference; • Software and its engineering;

## KEYWORDS

Systematic Literature Review, Language Models, Deep Learning, BERT

## ACM Reference Format:

Rand Alchokr, Manoj Borkar, Sharanya Thotadarya, Gunter Saake, and Thomas Leich. 2022. Supporting Systematic Literature Reviews Using Deep-Learning-Based Language Models. In *Proceedings of the 25th ACM/IEEE Joint Conference on Digital Libraries (JCDL '22)*, June 20–24, 2022, Cologne, Germany. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/000000.000000>

## 1 INTRODUCTION

Systematic Literature Reviews (SLRs) are a well-established research method used within the evidence-based research paradigm. In most scientific fields, such as medicine and software engineering, Systematic Literature Reviews have become a popular method for conducting literature analyses and meta-studies to aggregate evidence on a topic [17, 32]. Such types of studies are helpful to synthesize and assess the available empirical data, for example, of controlled experiments on the impact of identifier names. However, despite the methodological benefits of an Systematic Literature Review, it can be an extremely tedious and time-consuming process—especially with the number of scientific publications constantly increasing every year. This makes the task of identifying the most relevant publications an exceedingly difficult task for any researcher. Consequently, there has always been a need for automating Systematic Literature Reviews, and various tools provide different types of support [14, 22, 31]. Generally, tools assisting the conduct phase retrieve publications based on a query that uses *term frequency and inverse document frequency* (TF-IDF) to identify similar publications, and display these to the user. However, one of the main drawbacks of such existing tools is their inability to capture semantics of the actual texts, which could help to improve the correct selection of publications according to the specified query and publications previously selected by the user.

In this paper, we address the tedious process of identifying relevant primary studies during the conduct phase of an Systematic

Literature Review. For this purpose, we use deep-learning architectures in the form of the two language models BERT [7] and SBERT [30] to semi-automate this phase, and thus support the entire Systematic Literature Review method. We use the learned contextualized meanings of the titles and abstracts of publications to learn clusters and group publications accordingly. Learning contextualized relationships at word, sentence, or paragraph level is possible through the two language models. To evaluate our proposed technique, we compare the models' resulting clusters of publications with the outcomes of two Systematic Literature Reviews conducted manually by researchers (i.e., we compare the publications selected by the models and the researchers).

In detail, our contributions are:

- We propose a semi-automated technique to cluster the initial set of publications of an Systematic Literature Review using contextualized embeddings obtained from variants of the BERT language model.
- We introduce a weighting scheme to add value to publications close to the Systematic Literature Review search string.
- We evaluate our technique based on an empirical comparison with two existing manually conducted Systematic Literature Reviews.
- We publish a persistent open-access repository comprising the source code and datasets we used for the evaluation.<sup>1</sup>

Our results indicate that our proposed technique is a promising way to improve the automation of conducting Systematic Literature Reviews. We sketch our further plans of improving, integrating, and evaluating our technique in more detail, which highlights interesting research opportunities.

In section two, we discuss Systematic Literature Review methodology with relevant work of automation tools, an overview of language models and neural networks. We describe our contribution to the conducting phase of the Systematic Literature Review approach and the models adopted to generate the embeddings in section three. Next in section four, we discuss the evaluation of our experiments and its results. To our knowledge, this is one of the first works involving deep learning infrastructures with Systematic Literature Reviews. In section five, we list several threats to validity and finally we conclude in section six that our novel approach helps in reducing the time and effort required for identifying the primary studies within the conducting phase of Systematic Literature Reviews and encourages more intensive research towards integrating deep learning architectures in the form of language models with the Systematic Literature Review process.

## 2 BACKGROUND AND RELATED WORK

In this section, we discuss the Systematic Literature Review methodology and previous work related to its automation. We also, discuss some recently proposed techniques related to neural network and language models.

### 2.1 Systematic Literature Review

Typically any scientific research work begins with a literature review. However, a literature review is not considered of a significant scientific value if not conducted properly following specific guidelines. An Systematic Literature Review has three main phases, where each of them is divided into several steps, as proposed by Kitchenham and Charters [18], and summarized in Figure 1.

- (1) Planning: involves identifying the need for a review, establishing research questions and search queries, and developing review protocols.
- (2) Conducting: involves the retrieval of relevant research, selecting the primary studies, and data synthesis.
- (3) Reporting: involves formulation of the final report with the findings and analysis of the gathered evidence.

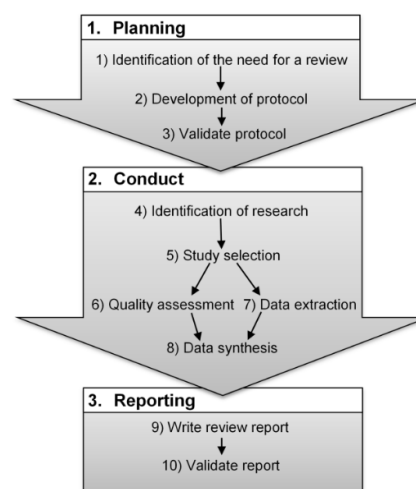


Figure 1: The three phases involved in a Systematic Literature Review process.

The goal of an Systematic Literature Review is to summarise existing evidence and identify gaps. However, the labor-intensive, error-prone, and time-consuming nature of this method has led researchers to develop a range of tools to provide support by automating parts of its process. (SLuRp) [3], (StArt) [15], (SLR-Tool) [11] are some of the very few tools that support all stages of the Systematic Literature Review. Although these tools support the entire Systematic Literature Review process, some of them have difficulties such as, complex installation procedures or lack of support for protocol development or collaboration. Some of the proposed tools majorly contribute to the conducting phase, where the main focus lies on primary studies' selection and data synthesis. While other tools, such as (PEx) [9] and (ReVis) [8, 10] employ text mining and visualization techniques to construct mappings between the research documents, site content analyzer tool uses entity recognition methodologies [5] and SLRONT [6] uses evidence-based ontology to support Systematic Literature Reviews.

Several studies evaluated some of the available tools based on a criteria of features [23, 24]. Their results show a need for more research employing new techniques to improve the effectiveness

<sup>1</sup>[https://www.dropbox.com/s/q8ong4217r1ioz9/DL\\_SLRAttachment.zip?dl=0](https://www.dropbox.com/s/q8ong4217r1ioz9/DL_SLRAttachment.zip?dl=0) (will be published on Zenodo)

and efficiency in identifying the relevant primary studies in the conducting phase. A required improvement would be capturing the semantic of the text. This learning of contextualized relationships at word, sentence, and paragraph level is made possible through the integration of deep learning architectures in the form of language models, and this is what we intend to achieve through our work.

## 2.2 Language Models

Over the last few years, language models, which are unsupervised multi-task learners, have become a crucial component of Natural Language Processing. They have been employed at the foundation of several Natural Language Processing tasks, such as, question answering, machine translation, reading comprehension, and summarization. These tasks which were typically approached with supervised learning on task-specific datasets are now tackled unsupervised by using language models. These models are pre-trained on gigantic amount of free-text data to learn predicting the probability of a sequence of words.

The earlier forms of language models used statistical techniques like linguistic rules, N-grams, and Hidden Markov Model (HMM) to learn the probability distribution over a sequence of words. Some count-based models such as Principal Component Analysis (PCA), and topic and neural probabilistic models [12, 27] involve unsupervised learning based on word frequency and global word co-occurrence matrix. The theory under study is that the words that occur together in the same context share similar and related semantic meanings. Context-based models, such as Skip-Gram [26] and Continuous Bag-of-Words (CBOW), model the text data using a sliding window of appropriate size that moves along the sentences and are trained to predict the probabilities of a target word/words given the rest of the context within the sliding window. Additionally, language models such as GloVe (Global Vectors) combine ideas of count and context-based models [13].

## 2.3 Neural Networks

Neural networks have shown important capabilities and effectiveness over the recent years in modeling language. Neural machine translation models were initially meant for language translation tasks [25]. They set the trend for using the word embeddings learned by the encoder of a bi-directional (LSTM) autoencoder for downstream tasks. Embeddings from language models (Elmo) picked this up a notch via producing contextualized word-level representations by pre-training the model in an unsupervised way [28]. Cross-View training is a semi-supervised learning approach where representations are improved by training in a supervised and unsupervised manner [4].

The current state-of-the-art language models utilize components of transformer architecture which inculcates multi-head self-attention mechanisms, thus helping the model learn long-range dependencies or relationships between words and sentences over longer sequences. Language models such as OpenAI-GPT [16], Bidirectional Encoder Representations from Transformers (BERT) [7], ALBERT [19], RoBERTa [21], and Text-to-Text Transfer Transformer (T5) [29] have brought the ImageNet moment in the field of Natural Language Processing. These models are trained on a giant collection of free text corpora and come in multiple variants

where contextual relationships and meanings are learned via techniques and tasks, such as masking, sentence order prediction, and translation.

Li et al. [20] employed a pre-trained BERT model to obtain sentence level embeddings and clusters on them to form sets of documents. While BERT has been trained with masked sequence and next sentence prediction tasks, Sentence-BERT (S-BERT) [30] is a modification of the BERT network to derive contextual and meaningful sentence embeddings using siamese and triplet networks. Additionally, Li et al. [20] used named entities as part of the weighting schemes to enhance the embeddings for clustering. We, on the other hand, explore clustering documents based on word, sentence, and paragraph-level embeddings and define our own weighting schemes based on the search string designed during the planning phase of the Systematic Literature Review process.

## 3 EXPERIMENTAL SETUP

In the following section, we explain our experiments beginning with identifying the dataset we use and the three basic modules we employ. These modules, as illustrated in Figure 2, involve:

- 1) Contextualized embeddings modules for getting various levels representations.
- 2) Weighting schemes module.
- 3) Clustering module.




Figure 2: Overview of our Framework.

### 3.1 Dataset

The dataset used in our experiments comprises of two collections, each belongs to an Systematic Literature Review study:

- SLR<sub>1</sub> "Bad Smell Detection Using Machine Learning Techniques\_a systematic literature review", 2019 [1]: 153 documents as an initial set of results from Scopus<sup>2</sup> and 12 out of them were manually identified as primary studies.
- SLR<sub>2</sub> "Automatic software refactoring\_a systematic literature review", 2019 [2]: 174 documents as an initial set of results from Scopus, out of which 9 are manually identified as primary studies.

For each Systematic Literature Review, we replicated the search methodology as their respective authors originally performed by applying the search string on Scopus and retrieving the initial set of results. The bibliographic information could be exported in .bib or .csv formats. This resulting dataset is considered our experiment's dataset. Each of the retrieved documents have entries, such as title, abstract, author, keywords, year of publication, publisher, venue, ISBN, and journal data.

As part of data gathering and cleaning, we first checked the dataset for duplicates, incomplete entries, and missing information.

<sup>2</sup><https://scopus.com>

Then we extracted the text entries only, namely: title, abstract, and keywords. We retain only the ones that have complete data for title and abstract, as they are essential for learning the document-level embeddings. However, incomplete and missing keyword entries are not considered as part of the retention criteria, because often a large number of research articles do not include keywords or index terms section, and eliminating such documents may lead to poor identification of relevant research.

Finally, after extracting and cleaning the data, each dataset contains four columns as illustrated in **Table 1** - **title**, **abstract**, **keywords**, and **cluster labels**. These cluster labels were manually added to indicate whether this document is among the primary studies set based on the decision of the researchers conducting the Systematic Literature Review. All documents are labeled as relevant (1), whilst the rest of the documents are assigned a non-relevant label (0).

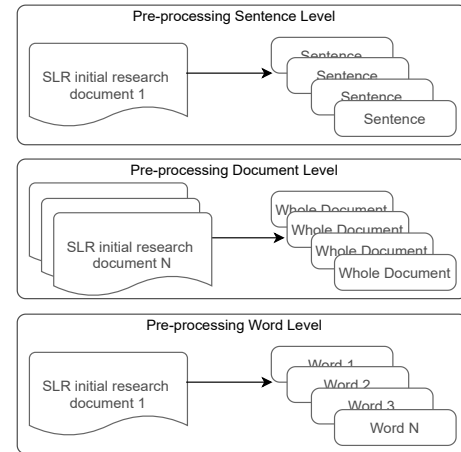
**Table 1: Overview of the document entries we extracted from Scopus.**

Variable name	Description
Title	The title of the document.
Abstract	The abstract of the document.
Keywords	The identified keywords of the document.
<i>Added manually:</i>	
Cluster label	A Boolean indicator→ primary study (1) or not (0).

### 3.2 Contextualized Word Embeddings - Documents Embedding Module:

Existing Systematic Literature Review tools retrieve documents based on the query string that uses (TF-IDF) where the similarity or dissimilarity of documents depends only on word frequencies. Inability to capture the semantics between text documents is one of the main drawbacks of the existing tools. To overcome this problem, here we propose to use contextualized word embeddings to capture the semantic similarity between documents. Contextual embeddings capture the usage of words across varying contexts and encode this information in a vector. This vector representation of a word changes when the context in which this word is used changes. We argue that using contextualized embeddings to represent documents will help end-users gain a better understanding of the underlying document space.

As we aim to understand how well the state-of-the-art deep learning architectures in the field of Natural Language Processing can help in capturing the intrinsic relationship between the documents. The documents embedding module forms the basis for learning clusters and grouping documents using the learned contextualized representations. In our experiment, we use BERT [7] and S-BERT [30] as our pre-trained language models to generate document-level contextualized embeddings. As mentioned before, BERT has been trained with masked sequence and next sentence prediction tasks. However, S-BERT is a modification of the BERT network to derive contextual and meaningful sentence embeddings using siamese and triplet networks. For implementation we used



**Figure 3: Pre-processing data for BERT, S-BERT models.**

Python and Scikit-learn<sup>3</sup> which is a free software machine learning library.

**3.2.1 Data pre-processing.** To explore what level of embeddings best captures the contextualized relationships between documents, we designed our experiments to accommodate learning of contextualized meanings at word, sentence, and document levels. In the case of BERT model, we design the input pipeline to obtain word and sentence level embeddings, while for S-BERT we modify the pipeline to obtain sentence and document level embeddings.

As shown in **Figure 3**, for word and sentence level embeddings, we break the document text into groups of individual words and sentences, respectively. However, at the document level, we treat the entire chunk of title and abstract of a document as one big sentence. Additionally, only in the case of BERT model, we tokenize the words and add special tokens such as [CLS], [SEP], and [PAD] to establish the start and end of sentences for the model to learn sentence level embeddings.

**3.2.2 BERT - Transformer based language model.** The architecture of BERT makes use of the transformer encoders to derive a language model trained on a humongous corpus. The usage of bi-directional multi-layer transformer encoders to capture the left-to-right context, as well as the right-to-left context, gives it the potential to capture long-term dependencies over past and future. To encourage the bi-directional nature, instead of training to predict the target given the context, BERT is trained on two tasks: *Masked Language Model* and *Next Sentence Prediction* - The former involves randomly masking a certain percentage of tokens in the sequence allowing the model to learn by predicting masked words without having prior information about the replaced words in the sequence. The second is inspired by the logical understanding that the downstream tasks rely more on sentence relationship than at word level and involve an auxiliary task of training a binary classification model to predict the next sentence given the previous sentence. BERT makes use of the same base model for all kinds of end-tasks, eliminating the idea of a task-specific model and improves upon

<sup>3</sup><https://scikit-learn.org/stable/>



the uni-directional nature of contextual understanding. It can be used for a large variety of Natural Language Processing tasks, often requiring a minimal amount of changes in the architecture, and have very few parameters to learn for each of the sub-tasks.

In our experiments with BERT for obtaining word-level embeddings, we encode the title of the documents and concatenate over the last four layers of hidden layer outputs to derive a fixed-sized embedding of 768 dimensions. For each document, we average over all the word embeddings for words occurring in the title of the documents. Similarly, for sentence-level embeddings, we break the abstract of documents into sentences and for each sentence, we derive a fixed 768 dimensions and average over all the sentence embeddings for each document to obtain document-level embeddings from them.

**3.2.3 Sentence-BERT (S-BERT).** S-BERT was fine-tuned on the Natural Language Inference (NLI) dataset using siamese and triplet network structures [30]. It was basically designed to perform Natural Language Processing tasks in low computational time. In S-BERT, a pooling operation is added to the output of BERT to generating a fixed-sized sentence embedding. S-BERT has three pooling strategies (i.e. [CLS] Token strategy, Mean strategy, and MAX strategy), evaluated on the STSbenchmark datasets<sup>4</sup>. The results showed that the Mean pooling strategy outperformed the other two strategies. Hence we have used it in our document clustering. The generation of document embeddings for a given document is carried out by a three-steps process. Figure 4 gives an overview of the implementation technique.

- 1) The text inside the document is split into sentences using the *Spacy sentence tokenizer* as sentences with a length greater than a certain threshold are removed. Based on the results of our experiments we set the threshold value to 20.
- 2) Each sentence is passed to S-BERT to derive a fixed-sized embedding of 768 dimensions.
- 3) Sentence embeddings of all the sentences belonging to a document are averaged to extract a document embedding.

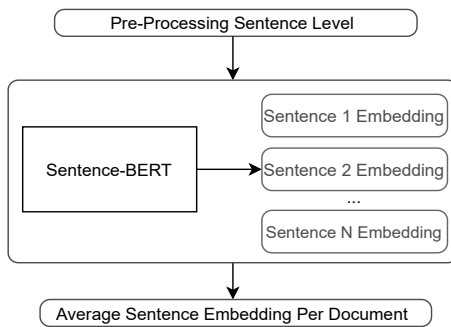


Figure 4: S-BERT implementation technique.

### 3.3 Weighting Schemes Module:

Systematic Literature Review mainly involves finding relevant research documents given a search string. The primary studies are

<sup>4</sup><http://ixa2.si.ehu.es/stswiki/index.php/STSbenchmark>

the relevant documents among the obtained initial set based on appropriate selection by the researchers. These are manually labeled as relevant label, meaning relevant for the study. Our weighting schemes is designed with the idea of recognizing relevant documents closer to the search string, unique words recognized by the highest (TF-IDF) scores of the entire set of documents and the search document.

The search string for finding primary studies in the conducting phase of the Systematic Literature Review contains keywords. We make use of this search string to identify the main keywords and create a list of study-relevant keywords. We also construct another keyword list, consisting of words that are identified by the highest (TF-IDF) scores. For every sentence in the document (D), a count  $N_D$  keeps track of the number of words from the search string and/or the number of words from the (TF-IDF) keyword list, that matches with the words in the sentences of (D). The document weight is:

$$W_D = \frac{N_D + 1}{N_S} \quad (1)$$

$N_D$  denotes the number of matching words in a sentence with the search string keyword list and/or (TF-IDF) keyword list.

$N_S$  denotes the total number of words in the sentence.

$W_D$  denotes the weighted contribution of each sentence to the document.

Figure 5 illustrates BERT, S-BERT weighting module.

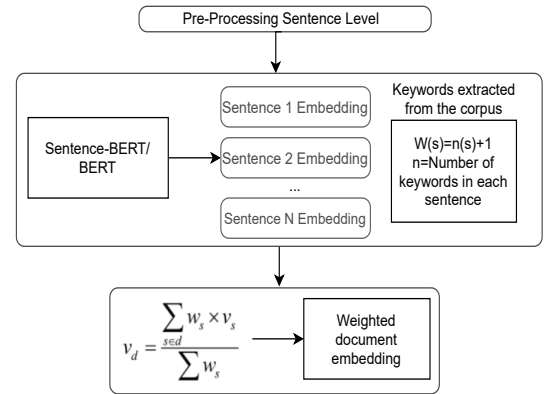


Figure 5: BERT, S-BERT Weightage module.

### 3.4 Document Clustering Module:

To recognize relevant and non-relevant documents during the conducting phase of the Systematic Literature Review, we use the clustering method. Clustering is a technique used in machine learning to group similar objects. The objects which are in the same cluster are more similar to each other as compared to the objects present in other different clusters. There are different kinds of clustering algorithms such as K-means clustering, Density-Based Spatial clustering, HDBSCAN, and Hierarchical clustering. In our experiments, we use *K-means* algorithm to group similar documents. The basic algorithm for K-means is described below in algorithm 1.

As an input to the clustering module, we use document level embeddings (obtained directly in case of paragraph-level or averaged over words and sentences in the document for word and

sentence-level structure respectively). We first identify the appropriate number of clusters for the embeddings using knee locator of *Kneedle algorithm* made available via Python kneed library<sup>5</sup>. This elbow determines the k value of the k-means algorithm to divide our embeddings into k clusters. Once k clusters are formed, we identify the cluster which consists majority of the documents labeled as primary study as per the true labels. This forms our relevant cluster.

---

**Algorithm 1** Basic Algorithm for K-means
 

---

**Input:** Document Embedding vectors and number of clusters  
i.e. K

**Output:** K clustered sets

- ```

0: procedure K-MEANS
1: Randomly select K centroids
2: Calculate the similarity of each Document Embedding with K centroids
3: Assign the Document Embedding vector to the cluster centroid whose distance from the cluster centroid is minimum of all the cluster centroids
4: Recompute the centroids of the newly formed clusters
5: repeat
   steps 3 and 4
6: until Centroids of newly formed clusters do not change
  
```
- 

## 4 RESULTS AND DISCUSSION

The designed experiments focus on exploring the possibilities over multiple configurations in different modules. Firstly, we experiment with the input data provided to determine how language models perform on different lengths and sorts of text. There are three ways in which we associate every document with a sequence of text.

- 1) Only the title is considered, we noticed that even though it has the shortest sequence length, it contains important keywords that define the document.
- 2) Only the abstract of each document represents the document as input data. Abstracts are slightly longer sequences of text in comparison to the title and often contain multiple sentences to form sentence-level relationships within the document set.
- 3) Both title and abstract are combined with subsequent word separators to form the biggest sequences of input data to be passed to the module.

In the document embeddings module, we use variants of BERT and S-BERT to compare how well they perform against each other in providing document-level embeddings from words and sentences respectively. Through the weighting schemes, we aim to find out if the identification of relevant documents for the search string or unique words of documents can become more prominent. We compare the clusters obtained using unweighted and weighted scheme approaches.

For evaluating the clusters on document embeddings, we use *Fowlkes-Mallows Index* (FMI), which is generally used when the ground truth labels are known. In our dataset, the labels are identified as relevant

<sup>5</sup><https://pypi.org/project/kneed/>

(1) and non-relevant (0) to the primary studies. FMI is defined as the geometric mean of the pairwise precision and recall:

$$FMI = \frac{TP}{\sqrt{(TP + FP)(TP + FN)}} \quad (2)$$

- TP : True Positives (i.e. number of pair of points that belongs to the same clusters in both true and predicted labels)
- FP : False Positives (i.e. the number of pair of points that belongs to the same clusters in true labels and not in the predicted ones)
- FN : False Negatives (i.e the number of pair of points that belongs in the same clusters in the predicted labels and not in the true labels)

We calculate the FMI scores for the entire document set as well as for the set manually identified as primary study. The score ranges from 0 to 1. A high value indicates a good similarity between the two clusters.

In our **first dataset SLR<sub>1</sub>**, 12 documents which are the primary studies form our relevant cluster. As noticed, S-BERT-weighted and S-BERT-paragraph achieve the best results in terms of 9 out of 12 documents being identified correctly as primary studies in the relevant cluster with a FMI score of 0.748 for the primary studies set. However, when we consider the number of additional documents in the relevant cluster, S-BERT-paragraph outperforms the S-BERT-Weighted.

For the **second dataset SLR<sub>2</sub>**, S-BERT-paragraph again achieves better results than the rest of the models in identifying 8 out of 9 primary study documents, but when we consider additional documents, the BERT-weighted approach performs better as it identifies 6 out of 9 primary studies and in addition has 18 other documents as part of the relevant cluster in comparison to S-BERT-paragraph which clusters 40 other documents in addition to 8 primary study ones.

Therefore, based on the performance, we can safely say that S-BERT-paragraph model fetches better results in terms of identifying more primary studies. However, when we compare the ratio of primary studies versus non-primary studies in the relevant cluster (i.e the number of additional documents identified), BERT-weighted is the best model with a consistent high FMI index score for each of the datasets.

Another interesting notion that comes into the picture when we have to make the model deployable is the execution time. This becomes important when the number of initial search documents extracted from Scopus is significantly larger. In terms of execution time, S-BERT-paragraph setting has lower values in comparison to other settings of the models.

## 5 THREATS TO VALIDITY

We replicated two Systematic Literature Reviews and compared our modules' selection and clustering of primary studies with the manual results achieved by the researchers. One potential threat could relate to deploying the search strings on digital libraries. The obtained results may change over a period of time resulting in differences in the datasets' results. Additionally, in our experiments, we made use of pre-trained models which have several advantages

**Table 2: Results for  $SLR_1$  - Total Documents = 153, Primary studies = 12.**

| Model              | Data             | # Clusters | Correct PS | MallowS for PS | MallowS all docs | # additional docs |
|--------------------|------------------|------------|------------|----------------|------------------|-------------------|
| TF-IDF             | Title + Abstract | 4          | 9/12       | <b>0.748</b>   |                  | 40                |
| S-BERT- WEIGHTED   | Title + Abstract | 4          | 9/12       | <b>0.748</b>   | 0.476            | 32                |
| S-BERT- UNWEIGHTED | Title + Abstract | 4          | 5/12       | 0.564          | 0.522            | <b>15</b>         |
| BERT- WEIGHTED     | Title + Abstract | 3          | 8/12       | 0.72           | <b>0.57</b>      | 22                |
| BERT- UNWEIGHTED   | Title + Abstract | 3          | 5/12       | 0.56           | 0.55             | 43                |
| S-BERT- PARAGRAPH  | Title + Abstract | 4          | 9/12       | <b>0.748</b>   | 0.49             | 22                |

PS = Primary studies, MallowS = Mallows score

**Table 3: Results for  $SLR_2$  - Total Documents = 174, Primary studies = 9.**

| Model              | Data             | # Clusters | Correct PS | MallowS for PS | MallowS all docs | # additional docs |
|--------------------|------------------|------------|------------|----------------|------------------|-------------------|
| TF-IDF             | Title + Abstract | 4          | 7/9        | 0.7637         |                  | 59                |
| S-BERT- WEIGHTED   | Title + Abstract | 5          | 7/9        | 0.7637         | 0.598            | 31                |
| S-BERT- UNWEIGHTED | Title + Abstract | 4          | 5/9        | 0.66           | 0.510            | 39                |
| BERT- WEIGHTED     | Title + Abstract | 3          | 6/9        | 0.66           | <b>0.61</b>      | <b>18</b>         |
| BERT- UNWEIGHTED   | Title + Abstract | 3          | 5/9        | 0.61           | 0.59             | 31                |
| S-BERT- PARAGRAPH  | Title + Abstract | 3          | 8/9        | <b>0.881</b>   | 0.56             | 40                |

PS = Primary studies, MallowS = Mallows score

in terms of execution time and ease of deployment and usage, however, when the number of initial documents is significantly larger, it makes more sense to fine-tune the model to the dataset and its vocabulary. Therefore, the evaluation so far needs more development, as it is infeasible to manually evaluate the method at large-scale. Adoption of analysis tools that can support the evaluation process, would considerably help research in this area.

## 6 CONCLUSION

In this paper, we present a novel approach to semi-automate the conducting phase of the Systematic Literature Review where identification and synthesis of relevant research called primary studies are of key interest. We simplify this process by integrating deep learning-based language models, specifically BERT and S-BERT in the setup. These models are utilized to extract contextualized embeddings from the title and abstract of documents at word, sentence, and paragraph levels. We also utilize a weightage scheme to give preference to documents that are in close relation with the Systematic Literature Review search string, and finally cluster the document level embeddings using the k-means algorithm to find clusters of similar documents.

To accomplish our tasks, we used two datasets each related to an Systematic Literature Review study. Our results demonstrate that clustering on embeddings at paragraph level using S-BERT-paragraph generally performs better than the rest of the model settings in terms of correctly identifying most of the primary study documents as relevant research. The results also outline that the weightage schemes consistently outperform unweighted approaches in cases of both BERT and S-BERT models.

To conclude, few of the miscellaneous experiments performed on our entire document set revealed promising results and could be used as a good starting point for identifying the primary studies for Systematic Literature Review. Learning contextualized relationships between documents using deep learning based architectures such as language models helped identifying most of the primary

studies and also helped in dealing with the drawbacks of the manual process of the Systematic Literature Review conducting phase generally.

## 7 FUTURE WORK

Engaging contextualized embeddings using deep learning based architectures such as language models in Systematic Literature Review is a promising aspect. More work is needed to explore the potential benefits of deep learning text analysis capabilities. This work paves way for ongoing research in this direction. An essential future experiment would be utilizing the full document text, instead of just the title and abstract. Another interesting area of exploration pertains to topic modeling experiments that we currently start investigating. Additionally, improving the evaluation part of the method by adopting analysis tools is an essential requirement. Therefore, to accelerate the process we invite interested researchers to join us in this research to provide a better level of reliability ending up with enhancing the conducting phase of Systematic Literature Review methodology.

## REFERENCES

- [1] Ahmed Al-Shaaby, Hamoud Aljamaan, and Mohammad Alshayeb. 2020. Bad Smell Detection Using Machine Learning Techniques: A Systematic Literature Review. *Arabian Journal for Science and Engineering* 45 (01 2020).
- [2] Abdulrahman Baqais and Mohammad Alshayeb. 2020. Automatic software refactoring: a systematic literature review. *Software Quality Journal* 28 (06 2020).
- [3] David Bowes, Tracy Hall, and Sarah Beecham. 2012. SLuRp - A tool to help large complex systematic literature reviews deliver valid and rigorous results. In *EAST'12 - Proc. 2nd Int. Work. Evidential Assess. Softw. Technol.* 33–36.
- [4] Kevin Clark, Minh Thang Luong, Christopher D. Manning, and Quoc V. Le. 2020. Semi-supervised sequence modeling with cross-view training. In *EMNLP 2018*. 1914–1925.
- [5] Daniela Cruzes, Manoel Mendonça, Victor Basili, Forrest Shull, and Mario Jino. 2007. Automated Information Extraction from Empirical Software Engineering Literature: Is that possible?. In *ESEM 2007*. 491–493.
- [6] Daniela Cruzes, Manoel Mendonça, Victor Basili, Forrest Shull, and Mario Jino. 2007. Using context distance measurement to analyze results across studies. In *ESEM 2007*. 235–244.
- [7] Jacob Devlin, Ming Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of deep bidirectional transformers for language understanding. In *NAACL HLT 2019*. 4171–4186.

- [8] Katia R. Felizardo, Gabriel F. Andery, Fernando V. Paulovich, Rosane Minghim, and José C. Maldonado. 2012. A visual analysis approach to validate the selection review of primary studies in systematic reviews. *Information Software Technology* 54, 10 (2012), 1079–1091.
- [9] Katia Romero Felizardo, Elisa Yumi Nakagawa, Daniel R.C. Feitosa, Rosane Minghim, and José Carlos Maldonado. 2010. An Approach Based on Visual Text Mining to Support Categorization and Classification in the Systematic Mapping. In *EASE 2010*.
- [10] Katia R. Felizardo, N. Salleh, Rafael M. Martins, Emilia Mendes, Stephen G. Macdonell, and José C. Maldonado. 2011. Using visual text mining to support the study selection activity in systematic literature reviews. In *ESEM 2011*. 77–86.
- [11] Ana M. Fernández-Sáez, Marcela Genero Bocco, and Francisco P. Romero. 2010. SLR-Tool a tool for performing systematic literature reviews. In *ICSOFT 2010*, Vol. 2. 157–166.
- [12] Michael Gutmann and Aapo Hyvärinen. 2010. Noise-contrastive estimation: A new estimation principle for unnormalized statistical models. *Journal of Machine Learning Research* 9 (2010), 297–304.
- [13] Hanson, Edward R. 1971. Musicassette Interchangeability. the Facts Behind the Facts. *Journal of the audio engineering society* 19, 5 (1971), 417–425.
- [14] Edgar Hassler, Jeffrey C. Carver, David Hale, and Ahmed Al-Zubidy. 2016. Identification of SLR Tool Needs - Results of a Community Workshop. *Information and Software Technology* 70 (2016), 122–129. <https://doi.org/10.1016/j.infsof.2015.10.011>
- [15] Elis Hernandes, Augusto Zamboni, Sandra Fabbri, and André Di Thommazo. 2012. Using GQM and TAM to evaluate StArt – a tool that supports Systematic Review. *CLEI Electron. J.* 15, 1 (2012).
- [16] Jeremy Howard and Sebastian Ruder. 2018. Universal language model fine-tuning for text classification. *ACL 2018* 1 (2018), 328–339.
- [17] Barbara A. Kitchenham, David Budgen, and O. Pearl Brereton. 2015. *Evidence-Based Software Engineering and Systematic Reviews*. CRC Press. <https://doi.org/10.1201/b19467>
- [18] Barbara A. Kitchenham and Stuart Charters. 2007. *Guidelines for Performing Systematic Literature Reviews in Software Engineering*. Technical Report EBSE-2007-01. Keele University and University of Durham.
- [19] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. 2020. Albert: A lite bert for self-supervised learning of language representations. In *ICLR 2020*. 1–17.
- [20] Yutong Li, Juanjuan Cai, and Jingling Wang. 2020. A Text Document Clustering Method Based on Weighted BERT Model. In *ITNEC 2020*. 1426–1430.
- [21] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. RoBERTa: A robustly optimized BERT pretraining approach. In *ICLR 2020*.
- [22] Christopher Marshall and O. Pearl Brereton. 2013. Tools to Support Systematic Literature Reviews in Software Engineering: A Mapping Study. In *International Symposium on Empirical Software Engineering and Measurement (ESEM)*. IEEE, 296–299. <https://doi.org/10.1109/esem.2013.32>
- [23] Christopher Marshall and Pearl Brereton. 2013. Tools to support systematic literature reviews in software engineering: A mapping study. In *ESEM 2013*. IEEE, 296–299.
- [24] Christopher Marshall, Pearl Brereton, and Barbara Kitchenham. 2014. Tools to support systematic reviews in software engineering: A feature analysis. In *EASE 2014*.
- [25] Bryan McCann, James Bradbury, Caiming Xiong, and Richard Socher. 2017. Learned in translation: Contextualized word vectors. In *NIPS 2017*, Vol. 2017-December. 6295–6306.
- [26] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. Efficient estimation of word representations in vector space. In *ICLR 2013*. 1–12.
- [27] Frederic Morin and Yoshua Bengio. 2005. Hierarchical probabilistic neural network language model. In *AISTATS 2005*.
- [28] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. Deep contextualized word representations. In *NAACL-HLT 2018*. 1–10.
- [29] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li Peter, and J. Liu. 2019. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research* (2019), 1–67.
- [30] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *EMNLP 2019*.
- [31] Yusra Shakeel, Jacob Krüger, Ivonne von Nostitz-Wallwitz, Gunter Saake, and Thomas Leich. 2020. Automated Selection and Quality Assessment of Primary Studies. *Journal of Data and Information Quality* 12, 1 (2020), 4:1–26. <https://doi.org/10.1145/3356901>
- [32] He Zhang and Muhammad A. Babar. 2013. Systematic Reviews in Software Engineering: An Empirical Investigation. *Information and Software Technology* 55, 7 (2013), 1341–1354.